

Streamer Channel Diagnosis Methodology v3

Advisor Review Request

Statistical and analytical methodology review — core diagnostic framework (§1–§8). Operational infrastructure, execution design, and portfolio bridge sections are excluded from this review scope.

Version: v3 draft
START · 2026-06-10

Review prepared:
2026-06-18

Scope: §0–§8 +
Appendices A/B/D/E/F

Pages excluded: §9–
§15, Appendix G–L

Reviewer Guide

WHAT THIS DOCUMENT IS

This is the core analytical methodology for diagnosing individual streamer channels on CHZZK (South Korea's primary live-streaming platform). It covers identity extraction, data collection, measurement correction, cohort construction, multi-axis diagnosis, tradeoff quantification, content channel analysis, and cross-validation.

The methodology has been applied to 2 cases so far (one LoL single-game non-virtual channel, one multi-game virtual channel). Both cases used cohorts of 130–320 channels with peak viewers in the 12–45 range.

Scope of this review: You are reviewing the *statistical and analytical methodology* — whether the diagnostic framework is sound, whether thresholds are justified, whether biases are accounted for, whether the cohort design is robust. You are NOT reviewing client management, deliverable design, execution coaching, or portfolio infrastructure — those sections are excluded.

What we want from you: Critique the methodology as if reviewing a pre-print's methods section. Flag unjustified assumptions, suggest alternative approaches, identify missing safeguards. If a threshold (e.g., 14h for long-

stream detection) looks arbitrary, say so and propose a data-driven alternative.

Table of Contents

- §0. Prerequisites (scope, data requirements, outputs)
- §1. Identity Extraction (3-layer validation)
- §2. Data Collection & Integrity Verification
- §3. Measurement Correction — Long-Stream Effects
- §4. Cohort Construction
- §5. 6-Stage Diagnostic Analysis
- §6. Target Tradeoff Quantification
- §7. Content Channel Effect Analysis
- §8. External Cross-Validation (XGPS)
- Appendix A. Risk Awareness Checklist (abridged)
- Appendix B/D. Trap Case Studies (anonymized)
- Appendix E. Cross-Case Comparison (anonymized)
- Appendix F. Data Schema Summary
- Review Questions (10 questions)**

§0. Prerequisites

0.1 Scope

This methodology applies to streamer channels meeting the following conditions:

- **Platform:** CHZZK (Twitch/AfreecaTV/YouTube Live applicable after cohort data availability check)
- **Scale:** 1,000–30,000 followers (above: cohort ceiling; below: data matching difficulty)
- **Data:** Access to 1 year of daily channel statistics
- **Category:** Categories where a cohort population of 200+ can be assembled

0.2 Required Data

Data	Minimum	Recommended	Ideal
Own daily stats	6 months	1 year	1 year+
Cohort population	Top 500 in category	2,000	Full
Follower data	For cohort matching	Top 6,000+	Full
Self-report	Channel intro / 1 interview	+ Viewer reaction sample	+ 1:1 interview
Content sample	1 recent broadcast	3-5 recent + shorts	Periodic samples
Collection provenance	source_url + date	raw snapshot + hash	Reproducible run package
Category basis	1 primary category	+ top category shares	90-day/1-year distribution
Peak basis	raw max	max + p95 + median + recent_median	Spike/regular tier separation

Below minimum = do not start. 6-month data misidentifies trends (see §3 traps). "No data" is classified as `absent` / `unknown` / `not_collected`, never treated as a conclusion.

0.3 Outputs

A single case produces 10 standard outputs. For this review, the relevant analytical outputs are: cohort data + methodology document (§4), 6-stage diagnosis summary (§5), tradeoff analysis (§6), content channel effect table (§7), and XGPS verification results (§8).

[Operational outputs — project registry, execution manual, tracking sheets, client deliverables — excluded from review scope]

§1. Identity Extraction

1.1 Input Sources

Three layers are collected in balance:

Layer	Source	Question
Self-perception	Channel intro, interview, self-statements	"How does the streamer define themselves?"
Content reality	Recent broadcast/short-form samples	"What are they actually doing?"
Viewer reception	Chat/comment/reaction patterns	"Why do viewers watch?"

1.2 Extraction Principle — 3-Layer Convergence

The *intersection* of all three layers is the true identity:

- Self says it but content doesn't show it → **Aspiration** (not identity)
- Content shows it but self is unaware → **Latent potential** (unconfirmed)
- Only viewers see it → **Misperception** (not identity)
- **All three present → Identity confirmed**

1.3 One-Line Positioning

The 3-layer convergence is compressed into a single sentence. Format varies by case — "*contrast appeal + relationship metaphor*" or "*intent-reception alignment*" — whichever captures the case's essence.

1.4 Entry/Retention Separation

On top of identity, two content functions are made explicit:

- **Entry (inflow) driver:** Why a viewer *first clicks* — surface-level, functional
- **Retention (stay) driver:** Why a viewer *stays* — emotional, relational

Confusing these two derails content strategy.

1.5 Vocabulary Handling

Streamer's own spontaneous vocabulary is treated as an *identity marker*. In analytical documents, it's preserved in quotes; in anonymized outputs, it's replaced with functional equivalents.

TRAPS

- Mistaking broadcast titles/surface expressions for identity without data confirmation.
- Accepting only self-report. Streamers tend to describe their *aspirational* self. Cross-validate with viewer reception and content reality.
- Confirming an integrated interpretation without the streamer's explicit validation.

§2. Data Collection & Integrity Verification

2.1 Channel Data

Collect 1 year of daily statistics. 6 months = do not start (see traps).

Required columns: date, followers (cumulative), broadcast hours, peak viewers, average viewers, chat (6-min average or cumulative), viewership (if available).

2.2 Cohort Population Data

Obtain **top 2,000 channel rankings** in the subject's category from external statistics services. Minimum columns: channel name, broadcast hours, peak viewers, average viewers, viewership.

2.3 Follower Matching Data

Cohort channels' follower counts may not appear in the population ranking. **Separately match follower rankings** (top 6,000+).

2.4 Verification Procedure

Check all 5 immediately upon receipt. **If any fails, do not proceed.**

1. **Period completeness** — No missing days? End date is current?
2. **Outliers** — Peak < average? Impossible values?
3. **Broadcast hour distribution** — Frequency of <0.1h or >24h entries?
4. **Category consistency** — Does the subject actually belong to this population's category?

5. **Cross-source match** — Same channel's numbers match between own data and population data?

2.5 Collection Provenance Schema

Every collected value must be traceable. Required metadata columns:

run_id, source_name, source_url, request_url, platform, cohort_type, aggregation_window_start, aggregation_window_end, fetched_at, http_status, content_type, raw_path, response_hash, parse_status, missing_reason, boundary_signal, collector_version

- Without raw_path , structured values are *notes*, not *canonical records*.
- Without missing_reason , blank cells are ambiguous between absence and parse failure.
- Without aggregation_window , ranking numbers risk cross-period contamination.

2.6 Absence Inventory

"Missing" is not discarded — it's stored as data in four states:

State	Meaning	Treatment
present	Confirmed to exist	Preserve source URL/raw path
absent	Checked, actually doesn't exist	Record absence source + confidence
unknown	Cannot determine yet	Leave open
not_collected	Not collected	Mark as out of collection scope

2.7 Privacy Minimization

Streamer diagnosis is public-channel-level analysis. No individual viewer/donor/account tracking. Chat/comments are anonymized by reaction type and vocabulary, never by individual user.

TRAPS

- **6-month trend judgment.** 6 months cannot control for seasonality. Always use 1 year.
- **Single-broadcast category assignment.** One recent broadcast $\neq$ category. Check distribution.
- **Population misidentification.** "Top 2,000 in category" $\neq$ "all streamers". No extrapolation.

§3. Measurement Correction — Long-Stream / No-Sign-Off Effects

3.1 Average Viewer Distortion Mechanism

Average viewers = total viewership $\div$ broadcast hours. When broadcast hours increase (including low-viewer late-night/dawn hours), the denominator inflates and the average is *mechanically* depressed. *This is a measurement effect, not a demand decrease.*

3.2 Long-Stream (No-Sign-Off) Identification Criteria

- **Threshold:** Daily broadcast time ≥ 14 hours
- **Habitual classification:** 5+ days/year at 14h+ OR par (broadcast hours / avg viewers) > 4

Days and channels meeting these criteria are flagged separately.

3.3 Correction Methods (choose 1 or combine)

Method	Calculation	Strength	Weakness
A. Peak ratio	Corrected avg = peak $\times 0.62$ (normal-channel median ratio)	Simple, long-stream robust	Vulnerable to single spikes
B. Time-weighted avg	$\Sigma(\text{avg} \times \text{hours}) \div \Sigma(\text{hours})$	Precise, strong for quarterly comparison	At channel level, denominator skews
C. Exclude long-stream days	Simple average excluding 14h+ days	Intuitive	Data loss

Recommended: Channel-level comparison (cohort) → A. Time series → B. Visualization → C.

3.4 Recommended Metric Hierarchy

Priority	Metric	Use
1	Peak viewers	Long-stream robust scale metric
2	Chat (6-min avg)	Core engagement
3	Corrected average viewers	Ratio calculations
Auxiliary	Raw average viewers	Never use alone

3.5 Zero Long-Stream Days Case

When the subject has 0 long-stream days in a year, §3 correction is *pass-through*. The output is the confirmation itself: "0 long-stream days confirmed, no correction applied." The metric hierarchy still applies; raw average can be used as-is.

TRAPS

- **Accepting a sudden avg-viewer drop as diagnosis.** Avg drop → first check broadcast-hour distribution for that period.
- **Uncritically accepting long-streaming as "strategy".** Whether it actually works must be measured.
- **Using only one corrected value.** Present both raw and corrected; verify the conclusion holds without correction.
- **Skipping §3 for zero-long-stream cases.** §3 is not "skipped" but "pass confirmed". Output must state "0 long-stream days confirmed."

§4. Cohort Construction

4.1 Peer Definition

Peak-based is default. Followers include cumulative/dormant accounts and misrepresent current live scale. Peak is a direct measurement of real-time scale.

- **Peer range:** Subject's peak × 0.5 to × 2.0

4.2 Filter Pipeline

Stage	Filter	Typical Size
Population	Top 2,000 in category	2,000
1st pass	Peak 0.5–2.0× filter	~200
2nd pass	Remove esports teams / tournament / corporate accounts	~190
3rd pass	Follower matching (separate data)	~150
Final	Remove clear outliers (dormant, <2 months active)	~130

4.3 Boundary Cases — Do Not Exclude by Surface

Boundary Case	Check Method	Decision
Virtual channels	Compare ratio/loyalty distributions vs non-virtual	If distributions overlap → include; clearly different → separate
Different game (secondary)	Category distribution (broadcast days)	If primary matches → include
Very small channels	Below peak floor	Separate pool

4.4 Robustness Verification (Required)

Before finalizing cohort conclusions, verify that **the subject's position holds across filter variants**:

Additional Filter	Purpose
Broadcast hours 10h+ / 20h+ / 30h+ only	Address short-broadcast ratio inflation
Single category only (subject's category)	Remove cross-game noise
Non-virtual only	Address virtual/non-virtual difference
Non-long-stream only	Check correction dependency

Subject holds top position across all filters → Robust. Any filter causes significant position shift → Investigate that factor deeply.

4.5 Category Confidence

Categories are not determined from a single recent broadcast. Use 90-day or 1-year category distribution when possible.

Confidence	Condition
high	Primary category share is dominant; profile/history/ranking align
medium	Primary visible but mixed; auxiliary cohort needed
low	Recent broadcast, profile, ranking disagree; manual review needed

4.6 Peak Basis Multiplexing

The 0.5–2.0× band uses peak as default, but `raw max peak` alone is vulnerable to raids/events/one-time spikes. Preserve multiple peak measures:

Basis	Use
<code>peak_max</code>	External ranking, maximum mobilization
<code>peak_p95</code>	Extreme spike mitigation
<code>peak_median</code>	Regular tier
<code>peak_recent_median</code>	Current tier
<code>peak_recent_p95</code>	Current tier including recent spikes

4.7 Primary / Auxiliary Cohort Operation

Boundary channels are not forced into a single cohort. **Primary cohort** = main-text conclusions. **Auxiliary cohort** = robustness verification. If the auxiliary cohort significantly changes conclusions, the report states "conclusion depends on market definition."

TRAPS

- **Excluding categories/virtual by intuition.** Check distributions first, then decide.
- **Population misidentification.** Is the population "LoL category" or "all CHZZK"? No extrapolation across boundaries.
- **Peak ceiling bias.** If peak 240 ceiling is used, large-channel ratios are lost. Verify conclusion impact.

§5. 6-Stage Diagnostic Analysis

5.1 Descriptive Statistics

Look at the *cohort distribution first*. Understand the population before locating a single channel.

- Ratio distribution (histogram, quantiles)
- Follower distribution
- Peak distribution
- Broadcast hour distribution

5.2 Multi-Axis Subject Positioning

At least 3 axes:

Axis	Meaning
Ratio (loyalty) rank	Efficiency — high engagement per follower?
Peak rank	Absolute scale
Follower rank (ascending)	Tier / weight class
Broadcast hour rank	Activity intensity

5.3 Homogeneity Check

Is the cohort actually a *homogeneous group*? Proxy measures for cohesion:

- **par (broadcast hours / avg viewers):** ≥ 4 = long-stream type, check distribution
- **Loyalty (peak / avg viewers):** 1.5–3 normal; above = one-time spike pattern

5.4 Key Relationship Quantification

Quantify the relationship between the subject's goal (e.g., 10K followers) and current strength (e.g., ratio) *using cohort data*:

- Follower vs ratio correlation (log-log)
- Median ratio by follower band
- Regression or quantile estimation

5.5 Correction Robustness

Check §3 correction dependency:

- Does the conclusion hold with raw metrics? (If not → conclusion depends on correction assumptions)
- What happens to the cohort when long-stream channels are removed?

5.6 Time Series Trend

Subject's 1-year quarterly trends using normal-day (non-long-stream) data:

- Quarterly peak viewers
- Quarterly chat
- Quarterly follower growth rate
- Relative speed of the three (which is fastest?)

5.7 Scheduling Effect Analysis

Performance is affected by broadcast time and day-of-week. Where possible, attach scheduling columns to daily data and analyze:

- Is growth from content change or scheduling change?
- Weekday/weekend and evening/late-night performance gaps?
- Do long broadcasts depress average viewers, or do they function as retention assets?

5.8 Spike vs Regular Tier Separation

Metric Movement	Interpretation
peak_max only rises	Event/raid/special content possible
peak_median rises	Regular-tier growth
avg_viewers rises	Retention/scheduling stability possible
chat rises	Core engagement strengthening
followers only rises	External inflow present, but live conversion may be weak

In §5 summaries, do not close with a single "growing/stagnating." Separate: **spike growth / regular-tier growth / follower growth / core-reaction growth.**

TRAPS

- Diagnosing a sudden avg-viewer drop before checking broadcast-hour distribution for that period.
- Claiming "growth" from peak_max alone when peak_median is flat.

§6. Target Tradeoff Quantification

6.1 Conflict Identification

Is there a conflict between the subject's *goal* (e.g., 10K followers, avg 100) and *current strength metric* (e.g., ratio)? Verify with cohort data.

6.2 Achievable Range Estimation

Estimate how far the strength metric drops when the goal is reached:

Method	Description
A. Cohort band median	Apply median of cohort's follower-band values
B. Log-log regression extrapolation	Extrapolate from fitted regression

When the two methods diverge significantly, present both as an *uncertainty range*.

6.3 Recommended Target + Alert Floor

- **Recommended target:** "Realistic upper tier" within cohort peers
- **Alert floor:** Below this → "Growth is not converting to genuine retention"

6.4 Perpetual-Intent Cases

When the streamer explicitly states *no intention to graduate* and *will continue even if shrinking*, the recommended target becomes a *waypoint* rather than a *destination*. Two additional questions:

1. **Threshold vs sub-threshold structure:** Aim for income threshold? Or build a sustainable sub-threshold structure?

2. **Deadline vs perpetual absorption:** Set a deadline for the target? Or absorb it as a waypoint in a perpetual flow?

Perpetual-intent is identified from interview signals: "no graduation plan," "will continue even if smaller," "no deadline."

6.5 Monetization Acceptance: Barriers Over Amounts

For revenue/donation analysis, record *whether the streamer can accept/request support* rather than exact amounts. Privacy-first: amounts are categorized, not stored precisely. Individual donor identification is never stored.

TRAPS

- **Assuming fixed-ratio growth.** "Maintaining current ratio at 10K" is nearly impossible. Acknowledge tier-up cost.
- **Binary framing.** Not "ratio vs scale, pick one" — the real question is "which ratio band to defend."
- **Ignoring perpetual intent.** Standard deadline-type targets can misalign with a perpetual-intent streamer.

§7. Content Channel Effect Analysis

7.1 External Content Conversion Measurement

Measure whether the subject's external channels (YouTube shorts, Instagram, Twitter, etc.) *convert* to main-channel growth.

Core question: Does reach (views) on external content convert to main-channel followers/viewers?

Method:

1. Identify external content posting dates
2. Compare main-channel follower growth rate in N days before vs after posting
3. Difference from baseline growth rate = conversion effect

7.2 Identity Alignment Assessment

Even if conversion happens, verify whether incoming viewers *match the streamer's identity*. Misaligned inflow doesn't retain.

Classification	Description
Aligned	Content matches main category
Partial	Identity adapted to trend
Misaligned	Pure trend-chasing (high reach, no identity connection)

7.3 Recommended Direction

- Aligned content → **Main lever**
- Partial → Supplementary, requires identity integration
- Misaligned → Non-priority (reach/refresh use only)

7.4 Content Segment Sampling

A broadcast is not one monolithic content piece. Actual appeal and conversion candidates emerge from specific segments within a broadcast. Sample segments with:

`content_mode, identity_marker, entry_hook, residue_hook, chat_spike, shorts_candidate`

7.5 External Channel Conversion Funnel

Separate reach from conversion. External content may have high views but zero main-channel inflow.

- **Reach metrics:** views, likes, comments, shares
- **Conversion metrics:** follower_delta, next_stream_peak_delta, next_stream_avg_delta
- **Alignment metrics:** identity_fit, CTA presence, main-channel link presence

7.6 Viewer Reception Vocabulary Sampling

Why viewers stay cannot be measured by numbers alone. From public reactions, record **reaction types and reception vocabulary** (anonymized — no individual tracking).

TRAPS

- **Judging effect by views alone.** Reach $\neq$ conversion. 50K views with zero main-channel inflow = weak effect.
- **Ignoring external channel as standalone asset.** Even without main-channel conversion, the external channel (e.g., YouTube subscribers) may be an asset on its own.

§8. External Cross-Validation (XGPS)

8.1 Validation Target Selection

Submit the *most reversible* or *most important* conclusions to external validation. Recommended targets (~3):

- The conclusion with the largest prior correction
- The core correction assumption (e.g., long-stream correction)
- The core strategic conclusion

8.2 Prompt Design — Hide Our Conclusions

The core XGPS principle: **prevent the external model from inferring our conclusion.**

Forbidden	Recommended
"Is our analysis that this channel is growing correct?"	"From the attached data, determine whether this channel is growing or stagnating."
"There's a long-stream artifact in this data, confirm it"	"Evaluate whether the average viewer metric can be used as the primary metric for this data."

Attach raw data only. Remove all our corrections and interpretations.

8.3 Convergence / Divergence Handling

Result	Interpretation
Convergence	Confidence increase — but if it's common-knowledge convergence, weak validation
Divergence	XGPS's real harvest. Possible blind spot. Re-verify with data.
Convergence with refinement	External provides sharper framing — accept it

Most common mistake: Reinterpreting divergence to defend our conclusion. Don't. Either settle it with data, or leave it as unresolved-open.

8.4 Same-Species Limitation

- If the external model is also an LLM, shared biases may exist
- Statistical-common-knowledge convergence is "common sense reconfirmation," not "independent validation"
- Real value is in *divergence* and *refinement*

8.5 Verification Input Package

Include	Exclude
Raw daily statistics	Our conclusion statements
Final cohort CSV	Correction interpretation statements
Cohort methodology summary	Client sensitive information
Content segment sample (partial)	Individually identifiable viewer reactions
Unresolved-open list	"Is this conclusion correct?" style leading questions

TRAPS

- **Defensive reinterpretation of divergence.** The impulse to rationalize unfavorable signals grows proportionally to their severity. Be aware.
- **Claiming convergence = "validation complete."** Convergence increases confidence but is not validation completion.

Appendix A. Risk Awareness Checklist (Abridged)

[Operational items A.6–A.9 removed from review scope]

A.1 Editor Overreach

- Did I assert a conclusion without team consensus?
- Did I designate a separate adjudicator for key judgments?
- Is quick consensus a signal that no real conflict was generated?

A.2 Self-Defense Bias

- Did I reinterpret external counter-evidence to defend our conclusion?
- Did I suppress unfavorable signals?
- Did I close something prematurely with "that's just [obvious reason]"?
- Did I assert an integrated interpretation without the subject's explicit confirmation?

A.3 Surface Judgment

- Did I exclude categories/types without data verification?
- Did I mistake broadcast titles/expressions for identity?
- Did I assert a trend from short-term (6-month) data?

A.4 Same-Species Dependency

- Did I over-interpret XGPS convergence as "independent validation"?
- Did I challenge the core conclusion with at least one different model?

A.5 Silent Domains

- Did I distinguish "truly absent" from "we didn't look"?
- Did I separate blank cells into absent / unknown / not_collected ?

- Did I record the original source of each absence (`absence_source`)?

Appendix B/D. Trap Case Studies (Anonymized)

Two cases applied so far. Case identifiers preserved; client identity removed.

Case 1: Single-Game, Non-Virtual

ID	Trap	Correction
K7	Nearly locked in a catchphrase-style broadcast title as identity	Demoted to "broadcast title" after data check; identity extracted separately
K8	Accepted a sudden avg-viewer drop as a real diagnostic finding	Corrected: measurement artifact from a 24h+ long-stream day
K11	Judged "stagnation" from 6-month data	1-year view showed trend reversal — short-term stagnation masking long-term rapid growth
O14	Nearly classified long-streaming as a "strategy"	Measurement showed it was non-recommended; reclassified
O16	Attempted to close XGPS divergence (engagement rate) as "scale effect"	Adjudicator flagged self-defense suspicion; retracted to unresolved-open

Case 2: Multi-Game, Virtual, Boundary Channel

ID	Trap	Correction
K1	Asserted "two-face integration" as identity from editor interpretation	Subject's explicit statement: "I perceive the two faces as separate" → partial separation
—	§1 extraction naturally triggered §6/§7/§9 statements	Separated §1 input from §6/§7/§9 entry data; linked to subsequent analysis
—	Jumped from §1 to §9 deliverable stage	Corrected by owner check; must pass §2–§8 first
—	Standardized subject's spontaneous vocabulary, killing the texture	Established vocabulary handling principle (§1.5)
—	Applied deadline-type standard target to perpetual-intent subject	Added perpetual-intent case handling (§ 6.4)
—	Zero long-stream days → tried to skip §3 entirely	Added zero-long-stream case handling (§3.5): "pass confirmed," not skipped

Appendix E. Cross-Case Comparison (Anonymized)

Axis	Case 1	Case 2
Category	Single game	Multi-game
Platform texture	Non-virtual	Virtual
Identity formation	Temperamental (natural expression)	Learned (multi-axis synthesis)
Self-doubt	Weak	Strong
Perpetual intent	Not stated	Strong
Long-stream	Correction applied	Correction not needed (zero days)
Data trend	1-year trend reversal (short-term stagnation → long-term rapid growth)	1-year gradual growth with possible spikes
Key finding	Diagnosis reversal via 1-year trend	Integrated low self-assessment diagnosis (pending validation)
One-line positioning format	Contrast appeal + relationship metaphor	Intent-reception alignment

Despite opposite starting textures, the methodology applied to both. With 3+ cases, type classification can be generalized.

Appendix F. Data Schema Summary

Full schemas are adapted per case. Core traceability columns (`source_*` , `raw_path` , `missing_reason` , `absence_source`) are preserved whenever possible.

F.1 Collection Provenance — `fetch_log.jsonl`

`run_id` , `source_name` , `source_url` , `request_url` , `platform` , `cohort_type` , `aggregation_window_start` , `aggregation_window_end` , `fetches_at` ,

http_status, content_type, raw_path, response_hash, parse_status,
missing_reason, boundary_signal, collector_version

F.2 Normalized Cohort — cohort_population_normalized.csv

run_id, cohort_type, source_name, channel_id, channel_name,
primary_category, category_basis, category_confidence, is_general_game,
is_virtual, total_stream_hours, stream_days, peak_max, peak_p95,
peak_median, peak_recent_median, avg_viewers, viewership,
follower_count, is_esports_team, is_tournament, is_corporate,
is_inactive, is_too_new, exclude_reason, collected_at

F.3 Final Cohort Derived — cohort_final.csv

channel_id, channel_name, final_cohort_group, peak_basis,
peak_ratio_to_target, within_peak_band, peak_to_followers,
avg_to_followers, viewership_per_hour, chat_per_avg_viewer,
peak_avg_ratio, par_hours_per_avg, nobangjong_flag, long_stream_flag,
final_include, final_exclude_reason

F.4 Content Segments — content_segment_sample.csv

vod_id, stream_date, segment_id, segment_start, segment_end,
content_mode, game_title, topic, identity_marker, entry_hook,
residue_hook, chat_spike, viewer_reaction_type, clip_candidate,
shorts_candidate

F.5 External Funnel — external_content_funnel.csv

platform, content_id, posted_at, content_type, identity_fit, views,
likes, comments, shares, cta_present, follower_delta_1d,
follower_delta_3d, follower_delta_7d, stream_peak_delta_next_stream,
conversion_assessment

F.6 Absence Inventory — absence_inventory.csv

absence_id, target_area, expected_asset, present_status, checked_source,
checked_at, absence_source, absence_type, confidence, impact,
next_check, notes

Review Questions

These are the specific points where we seek your critique. For each, we value: (a) whether the current approach is defensible, (b) specific

alternatives you'd recommend, (c) any biases or failure modes we haven't accounted for.

1 §3 — The 14-hour threshold for long-stream detection: is this cutoff justified, or is it arbitrary? Would you use a data-driven method (e.g., distribution inflection point, mixture model) instead?

Context: The 14h threshold was set empirically from Case 1, where a 24h+ broadcast day caused a visible avg-viewer artifact. The par > 4 secondary criterion was added as a ratio-based backstop.

HIGH If the threshold is wrong, measurement correction either misses distorted channels or over-corrects normal ones. Every downstream percentile is affected.

2 §3 — The peak ratio method uses 0.62 (normal-channel median ratio) as the correction factor. How robust is a single population median as a correction basis? Should channel-specific or band-specific ratios be used instead?

Context: 0.62 is the median of peak/avg ratio across non-long-stream channels in the cohort. It's applied uniformly to all corrected channels.

NORMAL Affects corrected avg-viewer values, but §3.5 robustness check verifies conclusions hold without correction.

3 §4 — The peak 0.5×-2.0× band for peer definition: is this range too wide, too narrow, or should it be adaptive (e.g., based on peak distribution density)?

Context: In practice, this produces ~200 channels from a 2,000 population. The band was chosen for balance between sample size and homogeneity.

HIGH Cohort definition is foundational. If the band systematically includes non-peers or excludes true peers, all §5 percentiles and §6 projections are biased.

4 §4 — Robustness verification uses 4 filter variants. Are these sufficient? Are there confounders (e.g., streaming schedule, content freshness, platform algorithm changes) that the current variants don't test?

Context: Current 4 variants test broadcast hours, category purity, virtual/non-virtual split, and long-stream inclusion. All applied to the same population.

NORMAL More variants increase confidence but also analysis cost. The question is whether 4 is the right cost-confidence tradeoff.

5 §4 — When the subject sits at the boundary between two cohorts (e.g., primary cohort percentile 72%, auxiliary cohort 45%), how should the gap be reported? Our current approach is to note it as a known limitation. Your earlier comment suggested follower-adjusted percentile or regression residual. We agree — which would you prioritize for this specific use case?

Context: This is the T1/T2 gap you flagged in the prior review. T1 (general game, n=184) and T2 (virtual, n=196) have different follower distributions, which we believe explains the gap but haven't formally proven.

HIGH If we can't explain the gap rigorously, the choice of "which cohort's percentile to trust" becomes a judgment call rather than a data-driven decision.

6 §5 — The efficiency percentile is adjusted via log-log regression ($\log(\text{efficiency}) \sim \log(\text{follower})$). In Case 1, this yielded $R^2 = 0.788$ (T1) and 0.749 (T2). Is the log-log specification appropriate here, or would you test alternatives (e.g., quantile regression, polynomial, GAM)?

Context: We chose log-log because efficiency = peak/follower follows a power-law relationship with follower count. The residual positions the subject relative to the regression line.

NORMAL The R^2 is high enough that specification choice may not matter much, but we want to know if we're leaving signal on the table.

7 §6 — Growth target estimation uses log-log regression extrapolation or cohort-band medians. Both extrapolate beyond the current cohort's range. What guardrails would you add to prevent over-extrapolation?

Context: The subject sits at ~700 followers with a growth target around 10K. The core cohort has peak 12-45, but 10K-follower channels are in the upper reference band (n=271), not the core cohort.

HIGH Over-extrapolation could produce unrealistic growth projections that damage client trust.

8 §7 — Conversion measurement uses a pre/post comparison (follower growth rate before vs after external content posting). What confounders should we control for? Would a difference-in-differences or synthetic control approach be more appropriate?

Context: Current method is simple: posting date identified → N-day pre vs post growth rate comparison. No seasonal or platform-level controls applied.

NORMAL The conversion measurement doesn't drive the core diagnosis but informs §10 execution recommendations (excluded from this review).

9 §8 — XGPS is designed to detect our blind spots by feeding raw data to external models without our conclusions. But if the external model is also an LLM, shared training biases may produce false convergence. How would you strengthen the cross-validation design?

Context: We acknowledge the same-species limitation (§8.4) and prioritize divergence over convergence. But we haven't implemented structural safeguards beyond awareness.

NORMAL XGPS is a supplementary check, not the primary validation. But its design determines whether it's genuinely useful or theater.

10 Overall — What statistical safeguards, biases, or failure modes are missing from this methodology that you would expect in a rigorous analytical framework? What's the single weakest link?

Context: The methodology was built incrementally across 2 cases. We suspect there are blind spots that only emerge at case 5+ or when applied to a different platform/scale.

CRITICAL This is the open-ended question. Everything else can be patched; structural blind spots can't be fixed if we don't see them.

NOTE TO REVIEWER

This methodology is a working draft (v3) with 2 completed cases. We are actively seeking critique before applying it to case 3+. Your review of the prior Part 1 channel diagnosis (T1/T2 percentile gap, regression residual suggestion) was exactly the kind of feedback we need — specific, method-level, with concrete alternatives. More of that, please.

Operational sections (§9–§15: deliverable design, execution coaching, portfolio bridge, disclosure management) are excluded from this review. If you have questions about how the analytical outputs are used downstream, we can discuss separately.

Document: methodology_review_danny_v3_20260618 · Prepared for advisor review · Sections §9–§15 and Appendices G–L excluded